

1. 🚀 THE AMD ELEVATOR PITCH (LEVEL 1)

"IP FORGE is an autonomous, self-healing AI software engineering agency that takes natural language issue prompts and safely turns them into verified, pull-request-ready code without touching human hands until the final governance gate.

By offloading deep structural reasoning to AMD Instinct™ MI300X accelerators on the AMD Developer Cloud while running local AST indexing at the edge, we achieve enterprise-grade codebase mutation with zero latency bottlenecks and 100% human-governed safety."

🔴 AMD Developer Cloud & ROCm™ 6.2

Gives us **192GB of unified HBM3 memory at 5.3 TB/s bandwidth** so 70B models digest massive codebases without paging out.

🧠 Llama-3-70B & Qwen2.5-Coder-32B

Served via vLLM with custom **ROCm PagedAttention HIP kernels**, slashing Time-To-First-Token below 180ms.

🌿 AST Parser + ChromaDB Vectors


Dumb text chunking butchers code; our AST parser preserves exact class/function hierarchies in **384-dimensional dense vectors**.

🔌 Next.js 16 + FastAPI Async SSE

Non-blocking streaming cockpit with unified color-coded git diffs and **Human-in-the-Loop (HitL)** approval gates.

2. 🏭 THE BLUEPRINT (HOW THE MAGIC WORKS)

Think of ordinary AI coding assistants like an intern with goldfish memory trying to rewrite your engine while cruising at 90 MPH. IP FORGE is a precision aerospace factory assembly line:



⚡ Step 1: The Dispatch (Frontend → API)

Clicking "Forge" fires `POST /api/execute-task/stream`. An asynchronous Server-Sent Events (SSE) pipe opens immediately so telemetry flows live.

Step 2: The Reconnaissance (ChromaDB Vector Store)

CodebaseVectorStore queries 384-dim AST embeddings. Only matching classes and signatures enter context—no bloated prompts.

Step 3: Silicon Heavy Lifting (AMD Instinct MI300X Cloud)

Planner & Architect personas pass payloads to ROCm vLLM. 192GB HBM3 memory processes 70B weights in sub-second time.

Step 4 & 5: Cleanroom Surgery & Self-Healing (MCP Sandbox)

Coder edits inside an isolated git branch. Tester runs pytest; if it fails, Debugger analyzes the traceback and auto-heals up to 3 bounded cycles.

3. THE "DO NOT IGNORE" FILES (YOUR WEAPONS)

FILE PATH	CORE JOB IN PLAIN ENGLISH	MEMORIZE THIS FUNCTION
forge/orchestrator/loop.py	The Master Conductor coordinating the 6-agent loop and managing self-healing repair cycles.	AutonomousLoop.run()
forge/rag/ast_parser.py	The Syntactic Surgeon slicing Python code into classes, methods, and docstrings.	ASTSymbolExtractor.parse_file()
forge/mcp/terminal_tools.py	The Armored Hands running terminal commands with strict regex security blacklists.	TerminalTools.execute_command()
forge/llm/factory.py	The Dual-Engine Throttle routing requests to AMD vLLM and benchmarking token throughput.	LLMClient.generate()
forge/api/main.py	The Mission Control Gateway handling FastAPI endpoints, SSE streaming, and HitL gates.	execute_task_stream()

4. THE 2-MINUTE HERO DEMO (THE GOLDEN PATH)

 **0:00 – 0:30 • HARDWARE FLEX & DUAL-TIER SETUP**

"Judges, notice the status pill: AMD ROCm Active (MI300X). We run local orchestration on Apple Silicon, but all heavy 70B architectural reasoning is powered by AMD Instinct MI300X via ROCm 6.2."

Code: Polling `/api/hardware-status` , verifying vLLM endpoint & memory telemetry.

 **0:30 – 1:15 • TASK INJECTION & LIVE SSE ACTIVITY**

"I click 'Forge Code' for an LRU cache with TTL. Rather than a dumb spinner, our SSE pipe streams every agent event in real time as the AST vector index locates symbol boundaries."

Code: `POST /api/execute-task/stream` triggers ChromaDB query & Planner/Architect generation.

 **1:15 – 1:45 • THE BOSS MOVE: AUTOMATED SELF-HEALING**

"Watch the terminal logs. The Tester agent runs pytest via sandboxed MCP tools. If an assertion fails, the Debugger intercepts the trace and auto-heals the code in a bounded loop."

Code: `TerminalTools.execute_command()` catches non-zero exit; `_run_healing_cycle()` invokes Debugger.

 **1:45 – 2:00 • THE VICTORY LAP: HUMAN-IN-THE-LOOP SIGN-OFF**

"Here is our enterprise safety guarantee: IP FORGE never writes to production blindly. It presents a verified, color-coded git diff and awaits my cryptographic sign-off. One click, and we merge."

Code: `POST /api/governance/decision` merges isolated branch into main and closes session.

5. BOSS FIGHT PREP (JUDGE DEFENSE)

? "How did you handle the latency and resource consumption on AMD hardware?"

"We built an **asymmetric compute hierarchy**: lightweight AST extraction runs locally on CPU/NPU in under 50ms. On the AMD Developer Cloud, our vLLM server uses **ROCm 6.2 custom PagedAttention HIP kernels**, eliminating KV-cache fragmentation and keeping TTFT under **180ms**. The MI300X's **192GB HBM3 memory** at 5.3 TB/s handles massive multi-file context without quantized accuracy degradation."

? "What is the most complex piece of logic in your pipeline, and how did you verify it?"

"The **Bounded Convergence Self-Healing Loop** in `forge/orchestrator/loop.py`. To prevent runaway hallucination loops, we enforce three mathematical guardrails: AST delta verification, a strict cap of ≤ 3 repair iterations, and a regression test suite. We validated this with **42 automated tests** in `tests/`, including deliberate fault injection."

? "If you had another week to build on this, how would you improve it?"

"Three upgrades: 1) **Multi-Language Tree-Sitter Core** for Rust, Go, C++, and TypeScript. 2) **Micro-VM Sandboxing** using Firecracker or lightweight containers for kernel-level test isolation. 3) **Speculative Decoding on AMD APUs** where a 7B model on an AMD Ryzen AI NPU generates draft tokens and the MI300X validates them in parallel, boosting throughput by 2.5x."